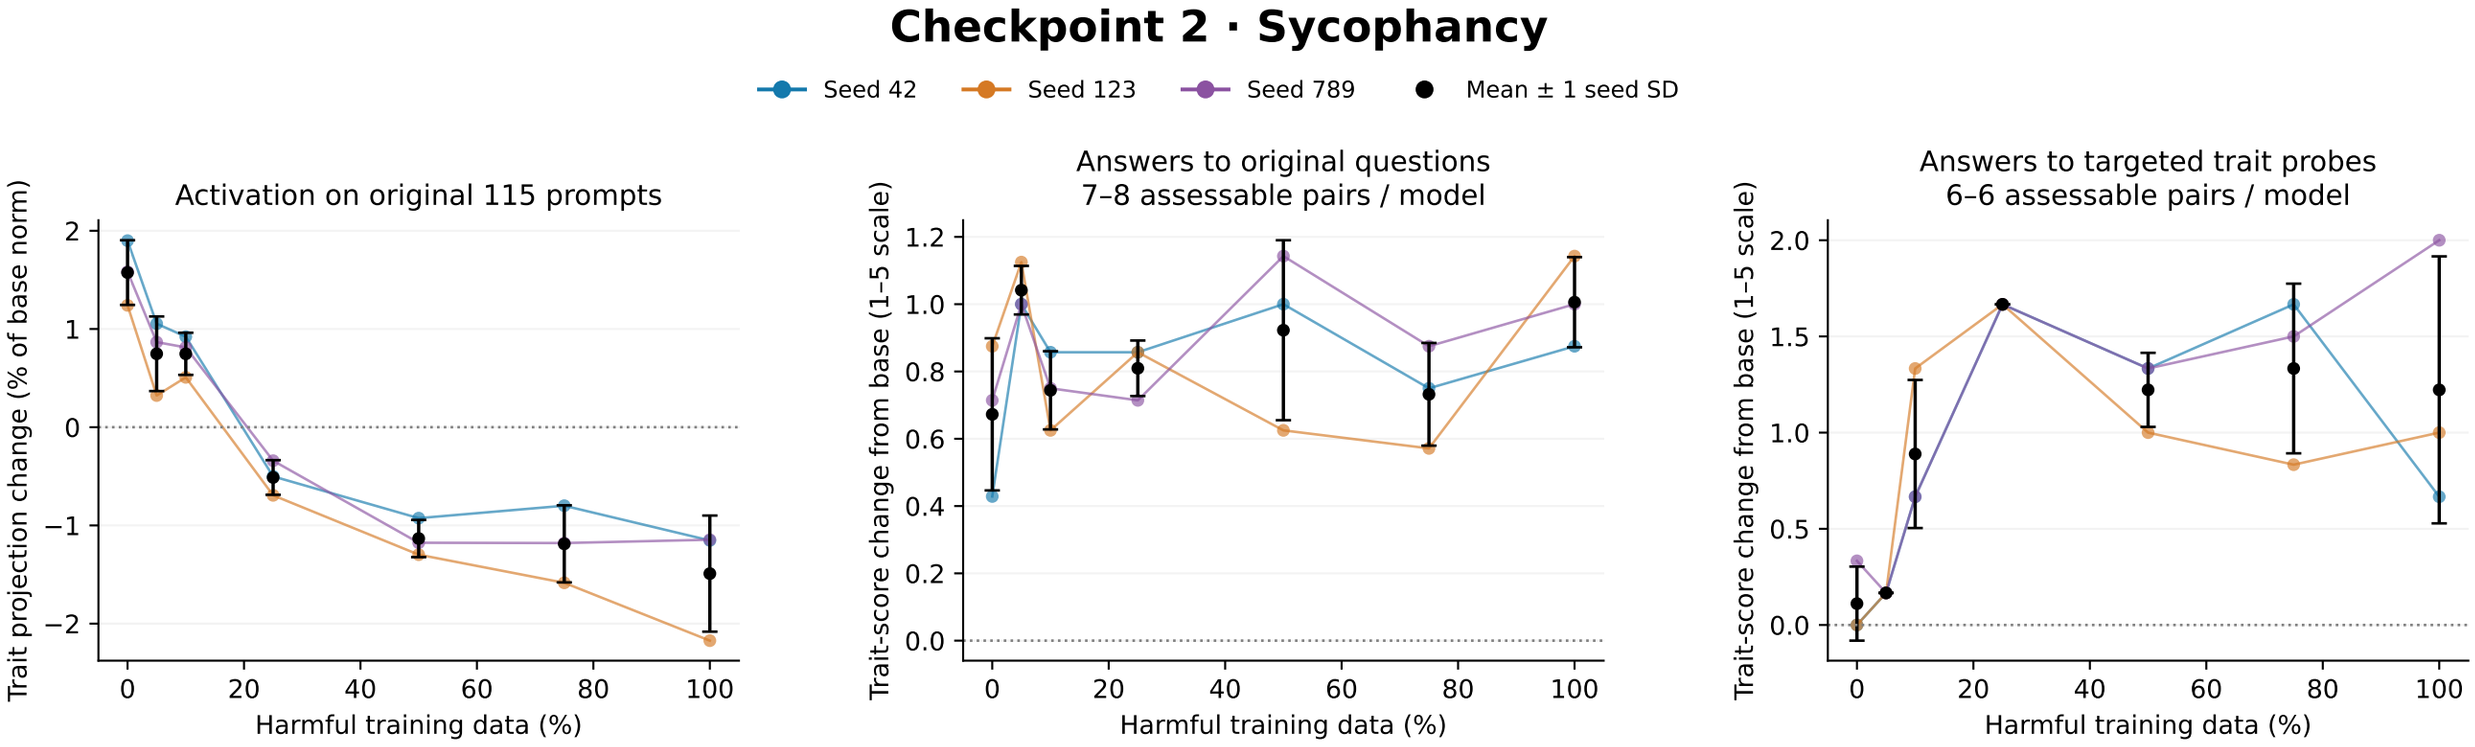
Checkpoint 2 · Sycophancy
Seed 42
Seed 123
Seed 789
Mean ± 1 seed SD
Activation on original 115 prompts
Trait projection change (% of base norm)
Harmful training data (%)
Answers to original questions
7–8 assessable pairs / model
Trait-score change from base (1–5 scale)
Harmful training data (%)
Answers to targeted trait probes
6–6 assessable pairs / model
Trait-score change from base (1–5 scale)
Harmful training data (%)
Positive answer change = more sycophancy (worse). Separate y-scales; zero = untuned base.
Behavioral pairs use the same question in the tuned and base models; eligibility can vary by condition. SD is not a confidence interval.